

# FIT1047 - Week 4

## Part 2: I/O and Interrupts



MONASH University

# Recap

We've now seen

- CPUs
- Memory

There's one component missing to complete the picture:

**Input and Output**

# Overview

Computers are useless without input/output.

Today:

- How does the CPU communicate with I/O devices?
- How does it handle time critical I/O?

# Early I/O

The first computers had very limited I/O.

- Punched paper tape or cards
- Teleprinters



Image source: Wikipedia

# Modern I/O

- Keyboard, mouse, touch pad, touch screen, voice control, gestures, accelerometers, barometers, GPS, ...
- Screens, printers, audio, smart home appliances, robots, ...

Also classed as I/O:


- Storage (hard disks, SSDs, SD cards)
- Network devices (WiFi, 4G, Ethernet)

# Modern I/O interfaces

I/O devices are now usually connected via *interfaces*:

- Standardised connectors and protocols
- Can be internal or external

Examples: USB, PCIe, HDMI, ...



# I/O and the CPU

The CPU needs to be able to

- Receive data from an I/O device
- Send data to an I/O device

Why send to input devices?

- E.g. set sensitivity, calibrate, switch on/off, ...

Why receive from output devices?

- E.g. check if ready, check if successful, ...

# I/O and the CPU

I/O devices have their own **registers**.

E.g. keyboard device: register holds the currently pressed key

How can we access these registers in machine code?

Two methods: memory-mapped and instruction-based I/O.



# I/O Access Method 1: **Memory-mapped**

Each I/O register is “**mapped**” to a **special memory address**. We can then use Load/Store.

**Example:** keyboard register is mapped to address A000.

- **Load A000** loads currently pressed key into AC register.

**Example:** printer register is mapped to address A100.

- **Store A100** stores current AC value in printer register (which then gets printed out).

# I/O Access Method 1: **Memory-mapped**

## Advantages:

- No need for new instructions
- Simple

## Disadvantages:

- We cannot use “mapped” addresses for memory any more
- The overall amount of usable memory is reduced (because some addresses are now unavailable)
- Programs may accidentally access I/O (when a program has a bug)

## I/O Access Method 2: **Instruction-based**

Add special I/O instructions to CPU ISA. For example, **Input** and **Output**.

Each I/O register still has an address (as in memory-mapped I/O).

But now these addresses are *separate* from the memory.

Example:

**Load A000** loads value from memory address A000

**Input A000** loads value from I/O register A000 (e.g. the keyboard)

# When to perform I/O

Now we know *how* to communicate with I/O devices.

But most I/O devices are *much, much* slower than the CPU.

So **when** should the CPU communicate with the device?

- How does it know that a key has been pressed?
- How does it know a printer is ready for receiving the next character?
- ...

Two methods: programmed and interrupt-based I/O.

# Programmed I/O

Programmer adds checks into code to periodically check I/O registers.

Also called **polling** I/O.

Pseudocode:

```
while (true) {  
    if (keyboardRegister.canRead()) {  
        processKeyboardRegister();  
    } else if (printerRegister.canWrite()) {  
        processPrinterRegister();  
    }  
}
```

# Programmed I/O

## Advantages:

- Very simple (no extra hardware needed)
- Programmer can decide how often to poll  
(e.g. 10000 times/second for network, 10 times/second for keyboard)

## Disadvantages:

- Programmer must be careful (poll registers regularly)
- Program is “I/O driven” (constructed around I/O)
- CPU is constantly in “busy loop”, causing high power usage

Programmed I/O is mostly used in embedded special-purpose systems.

# Interrupts

Opposite of polling:

- CPU is **notified** when it should communicate with I/O device
- CPU interrupts current program, executes special **interrupt handler** code, then continues normal program
- Programmer can write separate code for main program and interrupt handler (better separation of concerns)

# Interrupt signals

Device notifies CPU of pending interrupt by setting a bit in a special register.

CPU checks before each fetch-decode-execute cycle:

- Is interrupt bit set? Call interrupt handler.
- Otherwise: normal fetch-decode-execute.

Let's look at the RTL for this.



# RTL for Fetch with Interrupts

An interrupt handler is like a subroutine:

- Use **JnS** to
  - Store the **return address**
  - Jump to the code implementing the handler
- When the handler finishes, return to caller using **JumpI**

RTL for **JnS X**:

- |                            |   |                      |
|----------------------------|---|----------------------|
| 6. $MBR \leftarrow PC$     | } | Save current PC at X |
| 7. $M[MAR] \leftarrow MBR$ |   |                      |
| 8. $PC \leftarrow MAR$     | } | Jump to X+1          |
| 9. $PC \leftarrow PC + 1$  |   |                      |

# RTL for Fetch with Interrupts

1. If InterruptBit is 1:
    - 1a. Clear InterruptBit
    - 1b.  $MAR \leftarrow \text{InterruptHandler}$
    - 1c.  $MBR \leftarrow PC$
    - 1d.  $M[MAR] \leftarrow MBR$
    - 1e.  $PC \leftarrow MAR$
    - 1f.  $PC \leftarrow PC + 1$

}

Save current PC at **InterruptHandler**

}

Jump to **InterruptHandler+1**
  2.  $MAR \leftarrow PC$
  3.  $MBR \leftarrow M[MAR]$
  4.  $IR \leftarrow MBR$
  5.  $PC \leftarrow PC + 1$
  6. ...
- }

Normal fetch cycle

# Interrupt Handler

Must leave the CPU in the same state as before the interrupt!

- For MARIE: Contents of AC must be the same as before

This is called a **context switch**. It can be achieved by

- *Shadow registers:*  
When an interrupt happens, the CPU switches to a separate register file

Or

- Programming the interrupt handler to save registers to memory  
(see next example)

# Interrupt Vectors

How can the interrupt handler distinguish interrupts from different devices?

- Each device is assigned an **identification number**
- When raising an interrupt, the device **stores its number** in a special register
- The interrupt handler uses that number to jump to different subroutines
- The list of subroutines (one per device) is called an **interrupt vector**

# Example Interrupt Handler (MARIE syntax)

InterruptHandler, HEX 0

Store SaveAC

Save AC register to memory

Load InterruptVecAddr

Add InterruptDeviceID

Add device number to  
address of interrupt vector

Store Destination

JumpI Destination

Jump to device handler

SaveAC, HEX 0

/ Temporary storage for AC register

Destination, HEX 0

/ Computed destination handler

InterruptVecAddr, ADR InterruptVec

InterruptVec, ADR KeyboardHandler / device 0

ADR MouseHandler / device 1

ADR PrinterHandler / device 2

ADR SoundHandler / device 3

KeyboardHandler, ... / do some work

Load SaveAC

Restore AC register from memory

JumpI InterruptHandler


Return to normal program

# Interrupts in x86 PCs

Original design:

- 15 interrupt request (IRQ) signals
- hardware must be configured to use correct IRQ
- e.g. setting jumper on a network or sound card
- devices have to share IRQs

# Interrupts in x86 PCs



# Modern Interrupts

Use *Advanced Programmable Interrupt Controllers* (APICs).

- Integrated into CPUs
- Support more IRQs (results in fewer conflicts)
- Include high-resolution timers
  - E.g. cause an interrupt every millisecond
  - We will see later why that's useful!



# Disadvantages of interrupt-based I/O

- Different devices may need different priorities
  - E.g. keyboard (very low) vs graphics card (very high priority)
  - Can be achieved using different interrupt signals
- All memory transfers run through the CPU
  - E.g. read byte from disk, store into memory; or
  - Read byte from memory, transfer to graphics card
- I/O devices are fully controlled by CPU

Solution: DMA

# Direct Memory Access (DMA)

Modern CPUs can **delegate** memory transfer to dedicated controller.

- Hard disk controller can transfer data directly to RAM
- Graphics card can fetch image directly from RAM
- CPU is free to perform other tasks

CPU and DMA controller **share the data and address bus:**

- Only one can perform memory transfer at the same time

# Summary

## I/O

- Memory-mapped vs. instruction based
- Programmed (polling) vs interrupt-driven

## Interrupts

- Require *context switch* (call subroutine to handle interrupt, save registers)
- Jump into *interrupt vector* (one handler per device)

## DMA

- Off-load responsibility for data transfer to special controller
- Keeps CPU free to do useful tasks

# Outlook

Next lecture:

- Booting a computer
- BIOS / UEFI